

# Time Series Analysis

## Task: Choose a dataset with a time component and perform time series analysis

- Embark on a time series analysis project using a dataset with a time component, specifically historical stock prices. The objective is to uncover patterns, trends, and insights from the temporal data, enabling a better understanding of stock price movements over time

Time series analysis is a statistical method that involves studying and analyzing data points collected over time to identify patterns, trends, and make predictions about future values.

## Importing necessary libraries

```
In [1]: import pandas as pd
import numpy as np
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.tsa.seasonal import seasonal_decompose as sd
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```

## Loading and Exploring the dataset

```
In [3]: df = pd.read_csv('Microsoft_Stock.csv')
df.head()
```

```
Out[3]:
```

	Date	Open	High	Low	Close	Volume
0	4/1/2015 16:00:00	40.60	40.76	40.31	40.72	36865322
1	4/2/2015 16:00:00	40.66	40.74	40.12	40.29	37487476
2	4/6/2015 16:00:00	40.34	41.78	40.18	41.55	39223692
3	4/7/2015 16:00:00	41.61	41.91	41.31	41.53	28809375
4	4/8/2015 16:00:00	41.48	41.69	41.04	41.42	24753438

```
In [4]: df['Date'] = pd.to_datetime(df['Date'])
df.set_index('Date', inplace=True)
```

```
In [5]: df.sample(5)
```

```
Out[5]:
```

	Open	High	Low	Close	Volume
<b>Date</b>					
<b>2019-06-12 16:00:00</b>	131.40	131.97	130.71	131.49	17092464
<b>2018-01-05 16:00:00</b>	87.66	88.41	87.43	88.19	23407110
<b>2018-11-01 16:00:00</b>	107.05	107.32	105.53	105.92	33384201
<b>2020-03-03 16:00:00</b>	173.80	175.00	162.26	164.51	71677019
<b>2020-06-12 16:00:00</b>	190.54	191.72	185.18	187.74	43373587

```
In [6]: df.columns
```

```
Out[6]: Index(['Open', 'High', 'Low', 'Close', 'Volume'], dtype='object')
```

```
In [7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 1511 entries, 2015-04-01 16:00:00 to 2021-03-31 16:00:00
Data columns (total 5 columns):
#   Column  Non-Null Count  Dtype
---  ------  -
0   Open    1511 non-null     float64
1   High    1511 non-null     float64
2   Low     1511 non-null     float64
3   Close   1511 non-null     float64
4   Volume  1511 non-null     int64
dtypes: float64(4), int64(1)
memory usage: 70.8 KB
```

```
In [8]: df.describe()
```

```
Out[8]:
```

	Open	High	Low	Close	Volume
<b>count</b>	1511.000000	1511.000000	1511.000000	1511.000000	1.511000e+03
<b>mean</b>	107.385976	108.437472	106.294533	107.422091	3.019863e+07
<b>std</b>	56.691333	57.382276	55.977155	56.702299	1.425266e+07
<b>min</b>	40.340000	40.740000	39.720000	40.290000	1.016120e+05
<b>25%</b>	57.860000	58.060000	57.420000	57.855000	2.136213e+07
<b>50%</b>	93.990000	95.100000	92.920000	93.860000	2.662962e+07
<b>75%</b>	139.440000	140.325000	137.825000	138.965000	3.431962e+07
<b>max</b>	245.030000	246.130000	242.920000	244.990000	1.352271e+08


```
In [9]: df.isnull().sum()
```

```
Out[9]: Open      0
High      0
Low       0
Close     0
Volume    0
dtype: int64
```

## Visualization of Time Series

## (Exploratory Data Analysis)


```
In [10]: plt.figure(figsize=(10, 6))
plt.plot(df['Close'], label='Closing Price')
plt.title('Microsoft Stock Prices')
plt.xlabel('Date')
plt.ylabel('Stock Price')
plt.legend()
plt.show()
```



## Time Series Decomposition

Decompose the time series into its components: trend, seasonality, and residuals. This helps in understanding the underlying patterns.

```
In [11]: result = sd(df['Close'], model='additive', period=30)
result.plot()
plt.show()
```



## Statistical Analysis

- Conduct statistical tests for stationarity, such as the Augmented Dickey-Fuller (ADF) test.
- Check autocorrelation and partial autocorrelation functions to identify lag values for potential autoregressive (AR) and moving average (MA) terms in later modeling.


```
In [12]: adf_result = adfuller(df['Close'])
print('ADF Statistic:', adf_result[0])
print('p-value:', adf_result[1])
```

```
ADF Statistic: 1.7371362899270992
p-value: 0.9982158366942122
```

```
In [13]: adf_mean = df['Close'].rolling(12).mean()
adf_std = df['Close'].rolling(12).std()


plt.figure(figsize=(10,6))
plt.plot(df['Close'], color='blue', label='Original series')
plt.plot(adf_mean, color='yellow', label='Rolling mean')
plt.plot(adf_std, color='red', label='Rolling std. deviation')

plt.legend(loc='best')
plt.title('Rolling Mean and Std. Deviation')
plt.show()
```




## Plotting autocorrelation

```
In [14]: plot_acf(df['Close'], lags=30)
plt.show()
```



## Plotting Partial autocorrelation functions

```
In [15]: plot_pacf(df['Close'], lags=30, method='ywm')
plt.show()
```



## Resampling the data

(On a daily, weekly, and monthly basis and calculate the mean stock price)

```
In [16]: daily_prices = df['Close'].resample('D').mean()
print('Daily Variation in Stock Prices: ')
daily_prices.head()
```

Daily Variation in Stock Prices:

```
Out[16]: Date
2015-04-01    40.72
2015-04-02    40.29
2015-04-03     NaN
2015-04-04     NaN
2015-04-05     NaN
Freq: D, Name: Close, dtype: float64
```

```
In [17]: weekly_prices = df['Close'].resample('W').mean()
print('Weekly Variation in Stock Prices: ')
weekly_prices.head()
```

Weekly Variation in Stock Prices:

```
Out[17]: Date
2015-04-05    40.505
2015-04-12    41.540
2015-04-19    41.890
2015-04-26    43.950
2015-05-03    48.710
Freq: W-SUN, Name: Close, dtype: float64
```

```
In [18]: monthly_prices = df['Close'].resample('M').mean()
print('Monthly Variation in Stock Prices: ')
monthly_prices.head()
```

Monthly Variation in Stock Prices:

```
Out[18]: Date
2015-04-30    43.466667
2015-05-31    47.530000
2015-06-30    45.964091
2015-07-31    45.611818
2015-08-31    45.506667
Freq: M, Name: Close, dtype: float64
```

## Model Selection and Training (ARIMA)

- ARIMA stands for "AutoRegressive Integrated Moving Average"
- It combines three key components: AutoRegressive (AR), Integrated (I), and Moving Average (MA).
- ARIMA models are widely used for analyzing and forecasting time series data.


```
In [19]: df = df.reindex(pd.date_range(start=df.index.min(), end=df.index.max(), freq='D'))
```

```
In [20]: model = ARIMA(df['Close'], order=(5, 1, 5))
fit_model = model.fit()
```

```
C:\Users\suman\anaconda3\lib\site-packages\statsmodels\tsa\statespace\sarimax.py:9
66: UserWarning: Non-stationary starting autoregressive parameters found. Using ze
ros as starting parameters.
    warn('Non-stationary starting autoregressive parameters')
C:\Users\suman\anaconda3\lib\site-packages\statsmodels\tsa\statespace\sarimax.py:9
78: UserWarning: Non-invertible starting MA parameters found. Using zeros as start
ing parameters.
    warn('Non-invertible starting MA parameters found.')
C:\Users\suman\anaconda3\lib\site-packages\statsmodels\base\model.py:604: Converge
nceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
    warnings.warn("Maximum Likelihood optimization failed to ")
```

```
In [21]: # Visual Representation of AR model
plt.figure(figsize=(10,6))
plt.plot(df['Close'])
plt.plot(fit_model.fittedvalues)
plt.title('AR Model')
```

```
Out[21]: Text(0.5, 1.0, 'AR Model')
```




## Forecasting

```
In [22]: forecast = fit_model.get_forecast(steps=150)
forecast
```

```
Out[22]: <statsmodels.tsa.statespace.mlemodel.PredictionResultsWrapper at 0x2276292d210>
```

```
In [23]: # Visualization and Interpretation
plt.figure(figsize=(12, 6))
plt.plot(df['Close'], label='Historical Data')
plt.plot(forecast.predicted_mean, label='Forecast', color='red')
plt.title('Microsoft Stock Price Forecast')
plt.xlabel('Date')
plt.ylabel('Stock Price')
plt.legend()
plt.show()
```





```
In [24]: # Reporting
fit_model.summary()
```

Out[24]:

SARIMAX Results						
Dep. Variable:		Close		No. Observations:		2192
Model:		ARIMA(5, 1, 5)		Log Likelihood		-3328.254
Date:		Sun, 07 Jan 2024		AIC		6678.507
Time:		22:35:56		BIC		6741.120
Sample:		04-01-2015		HQIC		6701.392
		- 03-31-2021				
Covariance Type:		opg				
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.1114	0.054	-2.043	0.041	-0.218	-0.005
ar.L2	1.0703	0.014	76.145	0.000	1.043	1.098
ar.L3	-0.2797	0.057	-4.887	0.000	-0.392	-0.167
ar.L4	-0.9367	0.015	-61.010	0.000	-0.967	-0.907
ar.L5	0.0787	0.051	1.537	0.124	-0.022	0.179
ma.L1	-0.2097	0.049	-4.254	0.000	-0.306	-0.113
ma.L2	-1.0694	0.014	-78.950	0.000	-1.096	-1.043
ma.L3	0.5965	0.048	12.419	0.000	0.502	0.691
ma.L4	0.8781	0.019	45.140	0.000	0.840	0.916
ma.L5	-0.3739	0.045	-8.361	0.000	-0.462	-0.286
sigma2	4.1308	0.072	57.559	0.000	3.990	4.271
Ljung-Box (L1) (Q):		0.22	Jarque-Bera (JB):		23906.40	
Prob(Q):		0.64	Prob(JB):		0.00	
Heteroskedasticity (H):		17.54	Skew:		-0.55	
Prob(H) (two-sided):		0.00	Kurtosis:		19.14	

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [ ]:
```